

Parallel Programming in

OpenMP

– Labs –

OpenMP Exercises – I

- ❑ Write an OpenMP code to calculate π , using

$$\pi = \int_0^1 \frac{4}{(1+x^2)} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}$$

- ❑ implement the integrand as a function
- ❑ write your own reduction code
- ❑ use the OpenMP reduction clause
- ❑ compare the run-times

OpenMP Exercise – I

Starting point: numerical integration of $f(x)$

❑ smart OpenMP solution

```
int i, n;
double h, x, sum;

h = 1.0 / (double)n;
sum = 0.0;

#pragma omp parallel for default(none) \
    shared(n,h) private(i,x) \
    reduction(+: sum)
for(i=1; i<=n; i++) {
    x = h * ((double)i + 0.5);
    sum += f(x);
}
```

OpenMP Exercise – I

Sequential version: compiled with '-g -fast'

❑ runtime: 4.12 secs

```
int i;  
int N = 10000000000;  
double pi = 0;  
  
for( i=1 ; i<=N ; i++) {  
    pi += 4.0 / ( 1.0 +  
                ((i-0.5) / N) * ((i-0.5) / N) );  
}  
  
pi = pi * 1/N;
```

OpenMP Exercise – I

Automatic parallelization:

- ❑ compile with '-g -fast -xautopar -xreduction'

```
$ time OMP_NUM_THREADS=1 ./piauto
real    0m4.139s
user    0m4.125s
```

```
$ time OMP_NUM_THREADS=2 ./piauto
real    0m2.119s
user    0m4.211s
```

```
$ time OMP_NUM_THREADS=4 ./piauto
real    0m1.061s
user    0m4.153s
```

OpenMP Exercise – I

OpenMP version:

```
int i;
int N = 10000000000;
double pi = 0;

#pragma omp parallel for default(none) \
    shared(N) private(i) reduction(+: pi)
for( i=1 ; i<=N ; i++) {
    pi += 4.0 / ( 1.0 +
                ((i-0.5) / N) * ((i-0.5) / N) );
}

pi = pi * 1/N;
```

OpenMP Exercise – I

OpenMP parallelization:

- ❑ compile with '-g -fast -xopenmp'

```
$ time OMP_NUM_THREADS=1 ./piomp  
real    0m8.362s  
user    0m8.335s
```

But this is 2x slower!!!!

```
$ time OMP_NUM_THREADS=2 ./piomp  
real    0m4.256s  
user    0m8.404s
```

It scales!

```
$ time OMP_NUM_THREADS=4 ./piomp  
real    0m2.125s  
user    0m8.371s
```

What's going on here???

OpenMP Exercise – I

Looking at the problem:

- ❑ compiler comments:
 - ❑ almost no difference – same optimizations applied
- ❑ replaced loop body by a function call:
 - ❑ $\text{pi} += f(x)$
 - ❑ no effect!
- ❑ hmmm ... what now?

OpenMP Exercise – I

- ❑ What happens in the OpenMP version?
 - ❑ the code block following the “#pragma omp ...” gets “outlined” into a function
 - ❑ the compiler optimization is applied to this new, outlined function
 - ❑ in this process, some “information/knowledge” gets lost, and thus the compiler cannot apply advanced optimizations
 - ❑ in this case: the compiler does not optimize the division by N in the OpenMP version!

OpenMP Exercise – I

OpenMP version: solution

```
int i;
int N = 10000000000;
double pi = 0;
double h = 1.0 / N;

#pragma omp parallel for default(none) \
        shared(N,h) private(i) reduction(+:pi)
for( i=1 ; i<=N ; i++) {
    pi += 4.0 / ( 1.0 +
                ((i-0.5) * h) * ((i-0.5) * h) );
}

pi = pi * h;
```

OpenMP Exercise – I

OpenMP parallelization - fixed:

- ❑ compile with '-g -fast -xopenmp'

```
$ time OMP_NUM_THREADS=1 ./piomfix
real    0m4.137s
user    0m4.120s
```

```
$ time OMP_NUM_THREADS=2 ./piomfix
real    0m2.124s
user    0m4.177s
```

```
$ time OMP_NUM_THREADS=4 ./piomfix
real    0m1.061s
user    0m4.174s
```

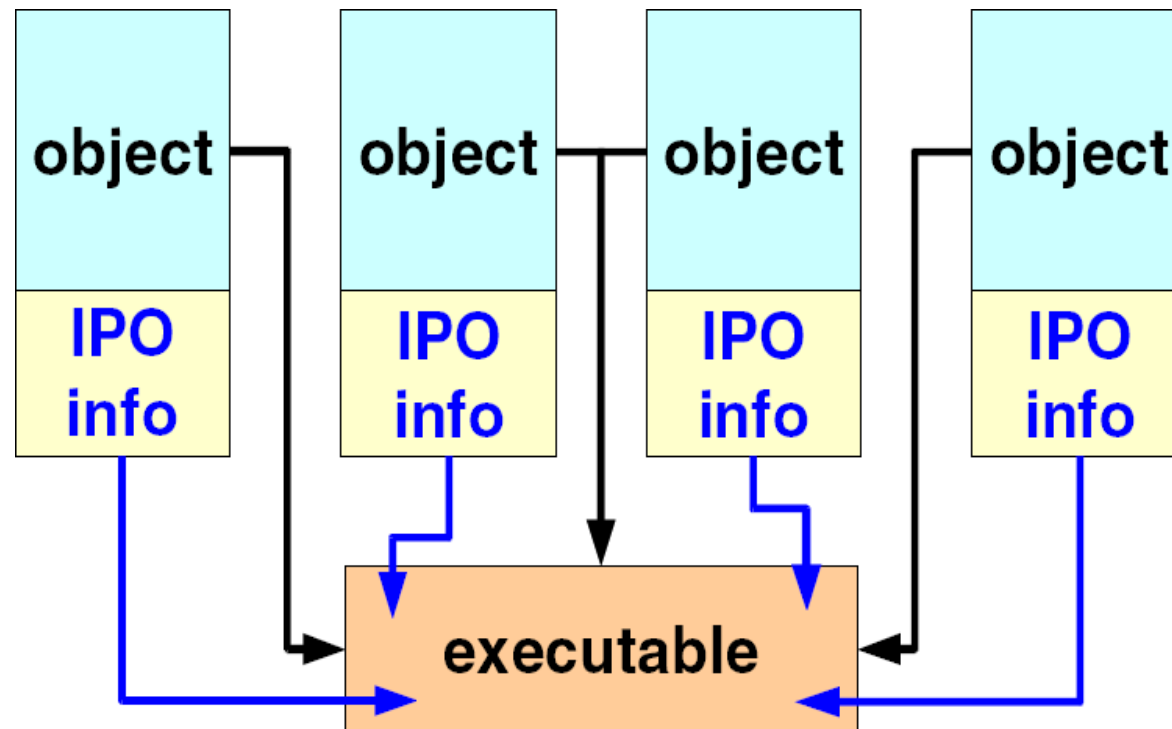
Same as with 'autopar'!

OpenMP Exercise – I

- ❑ Using more than one source file
 - ❑ if the function $f()$ that should be integrated is not in the same source file, automatic parallelization does not work
 - ❑ workaround: use IPO (Interprocedural Optimization)
 - ❑ with the Studio compilers: `-xipo`

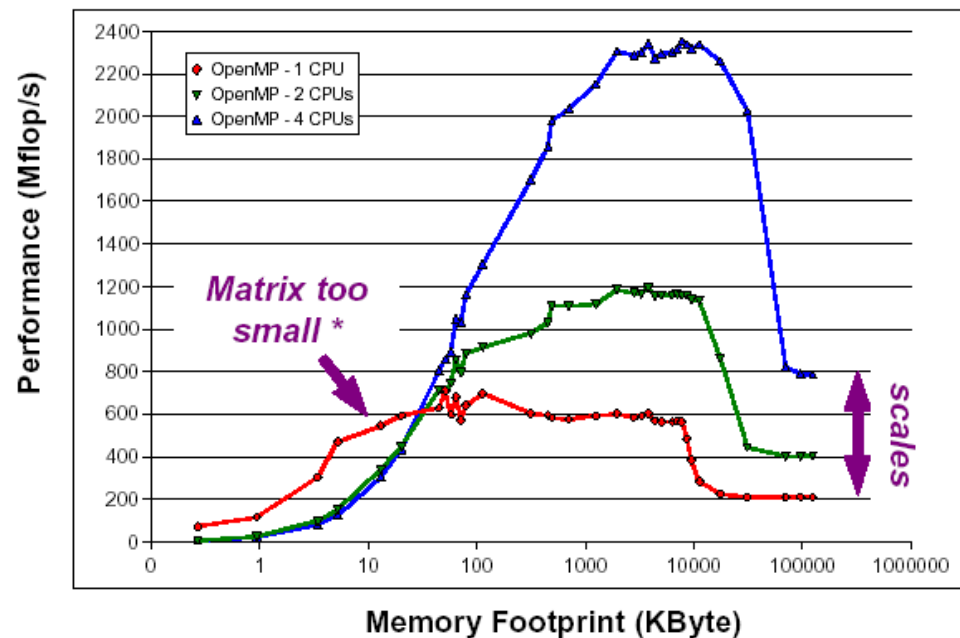
Re-cap: Inter Procedural Optimization

- ❑ With the -xipo option, the compiler stores additional information into the object files
- ❑ This information is used during the link phase to perform additional optimizations



OpenMP Exercises – II

- ❑ Improve the matrix times vector example by adding an if-clause to the omp pragma – experiment with the threshold value!



SunFire 6800
UltraSPARC III Cu @ 900 MHz
8 MB L2-cache

*) With the IF-clause in OpenMP this performance degradation can be avoided

Automatic Parallelization

```
1 void
2 mxv(int m,int n,double *a,double *b,double *c) {
3
4     int i, j;
5     double sum;
6
7     for (i=0; i<m; i++) {
8         sum = 0.0;
9         for (j=0; j<n; j++)
10             sum += b[i*n+j]*c[j];
11         a[i] = sum;
12     }
13 }
```

cc -g -fast -xrestrict -xautopar -xloopinfo -c mxv.c
"mxv.c", line 7: PARALLELIZED, and serial version
generated

"mxv.c", line 9: not parallelized, unsafe dependence
(sum)

Automatic Parallelization

❑ Compiling with '-xautopar':

```
void mxv(int m,int n,double *a,double *b,double *c) {  
    //... lines omitted ...  
}  
  
int main(...) {  
    //... lines omitted ...  
    // do max_it iteration for timing  
    for(int k=0; k < max_it; k++) {  
        mxv(m, n, a, b, c);  
    }  
    // ... lines omitted ...  
}
```

❑ compiler reports “PARALLELIZED” in mxv()

❑ but we see no speed-up!?!?

Automatic Parallelization

- ❑ Reason: the compiler “knows too much”!
- ❑ The compiler inlines `mxv()` in `main()`, and then there is a triple for-loop nesting, which prevents automatic parallelization! This can be seen in the compiler messages, too!
- ❑ It did create a parallel version, too, but does not use it!
- ❑ Fix: use `'-xinline=no'` => expected speed-up!
- ❑ Better: keep the source for `main()` and your subroutines in different source files!

OpenMP Parallelization

```
1 void
2 mxv(int m,int n,double *a,double *b,double *c) {
3
4     int i, j;
5     double sum;
6     #pragma omp parallel for private(i,j,sum)
7     for (i=0; i<m; i++) {
8         sum = 0.0;
9         for (j=0; j<n; j++)
10             sum += b[i*n+j]*c[j];
11         a[i] = sum;
12     }
13 }
```

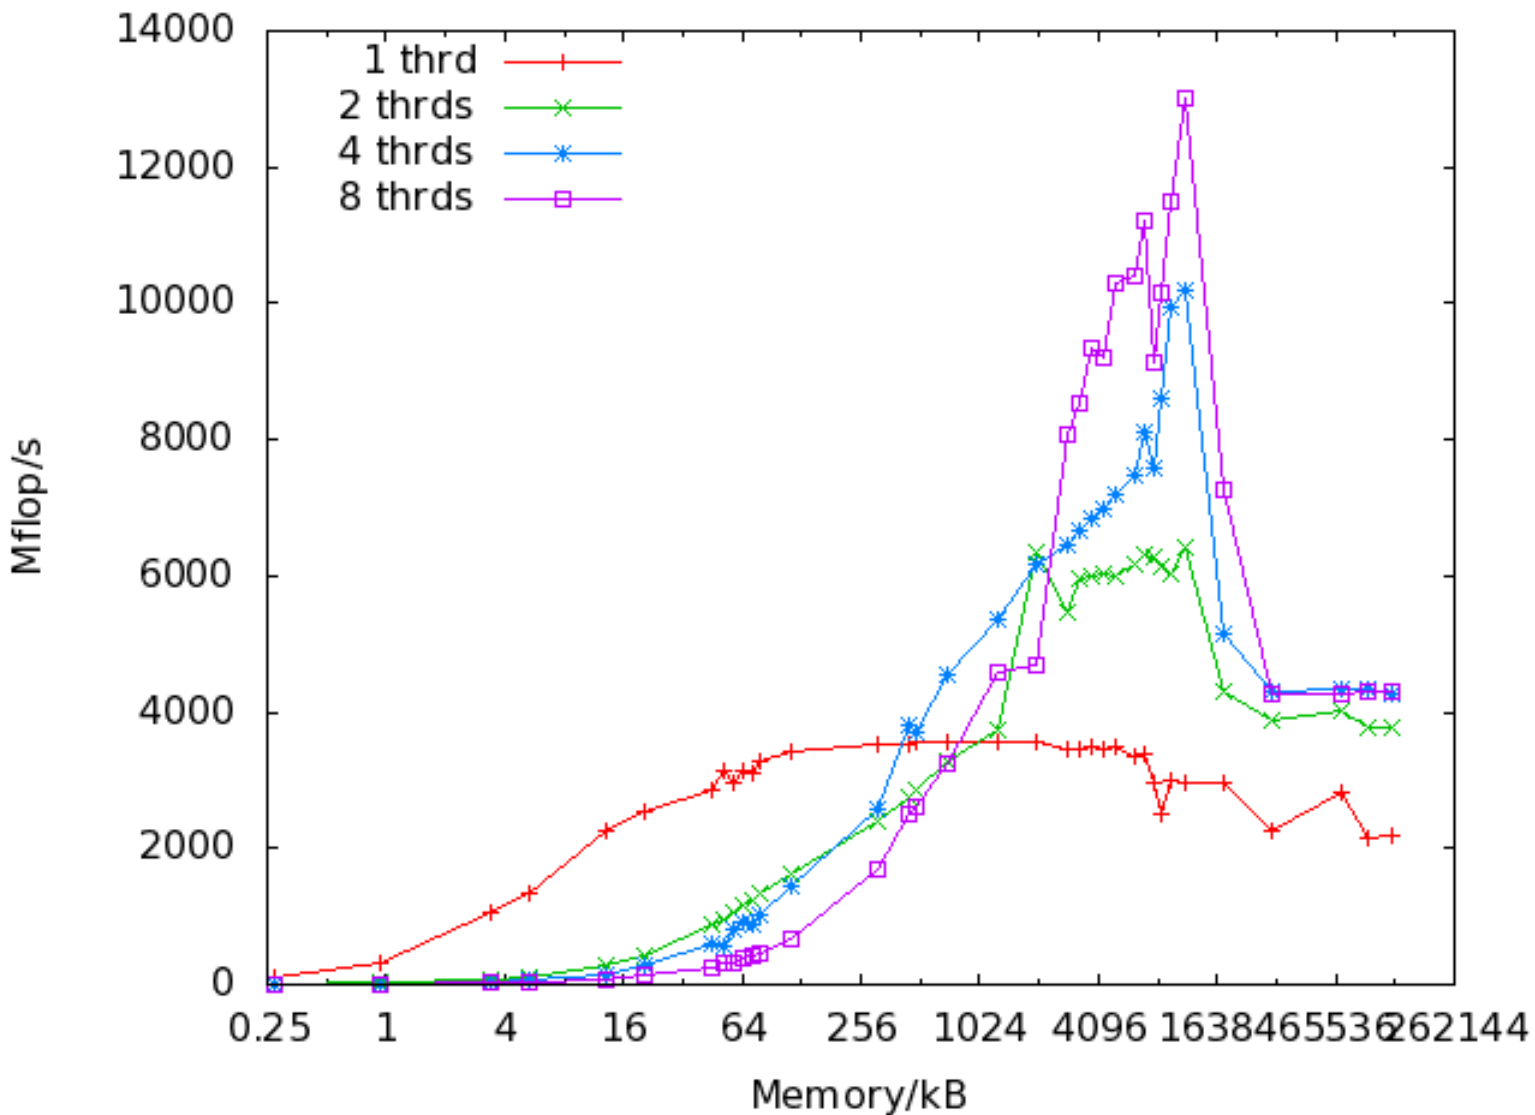
```
cc -g -fast -xopenmp -xloopinfo -c mxv.c
```

```
"mxv.c", line 7: PARALLELIZED, user pragma used
```

```
"mxv.c", line 9: not parallelized, loop inside
```

```
OpenMP region
```

OpenMP Parallelization



OpenMP Parallelization - if-clause

