

Parallel Programming in OpenMP – part II

Outline

- ❑ Data scoping – cont'd
- ❑ Orphaning
- ❑ Tasking
- ❑ OpenMP correctness & Data Races
- ❑ Runtime library
- ❑ Scheduling
- ❑ A real world example

OpenMP Syntax

More on data scoping



OpenMP Syntax

Reminder: the “private” clause –

- ❑ declares variables private to each thread:
`#pragma omp directive private (list)`
- ❑ i.e. a **new** variable is declared once for each thread
- ❑ all references are replaced with references to the newly declared variable
- ❑ variables declared private are uninitialized for each thread!



OpenMP Syntax

Consequences of private(...):

```
main() {
    ...
    A = 10;

    #pragma omp parallel
    {
        #pragma omp for private(i, A, B) ...
        for(i = 0; i < n; i++) {
            ...
            B = A + i;    // A undefined!
                        // unless declared firstprivate

            ...
        } /* end of omp for */

        ...
    } /* end of omp parallel */

    C = B;               // B undefined!
                        // unless declared lastprivate
}
```

OpenMP Syntax

Solutions:

```
#pragma omp ... firstprivate(list)
```

- ❑ All variables in list are initialized with the value the original object had before entering the parallel construct.

```
#pragma omp ... lastprivate(list)
```

- ❑ The thread that executes the sequentially last iteration updates all variables in list.

OpenMP Syntax

The “threadprivate” and “copyin” clauses:

- ❑ `threadprivate(list)`: creates a private copy of global data (e.g. common blocks or global variables in modules in Fortran) for each thread
- ❑ `copyin(list)`: copies the values from the master thread into the private copies
- ❑ subsequent modifications of list affect only the private copies – within one parallel region

OpenMP Syntax

Example 1:

```
int counter = 0;
#pragma omp threadprivate(counter)

int increment_counter()
{
    counter++;
    return(counter);
}
```

```
INTEGER FUNCTION INCREMENT_COUNTER()
COMMON/A22_COMMON/COUNTER
!$OMP THREADPRIVATE (/A22_COMMON/)
COUNTER = COUNTER + 1
INCREMENT_COUNTER = COUNTER
RETURN
END FUNCTION INCREMENT_COUNTER
```

OpenMP Syntax

Example 2:

```
int
increment_counter()
{
    static int counter = 0;
    #pragma omp threadprivate(counter)

    counter++;
    return(counter);
}
```

OpenMP Syntax

The `copyprivate(...)` clause

- ❑ copying a value out of a single region into the private data of other threads

```
#pragma omp single copyprivate(list)
{
    ...
}
```

```
!$OMP SINGLE ....
```

```
...
!$OMP END SINGLE COPYPRIVATE(LIST)
```

OpenMP Syntax

Example:

```
int x, y; /* global data */
#pragma omp threadprivate(x, y)

void use_values(int id, int a, int b) {
    printf("  TID %d: a = %d, b = %d, c = %d, d = %d\n",
           id, a, b, x, y);
}

void init(int id, int *a, int *b) {
    int r_a, r_b;

    #pragma omp single copyprivate(r_a, r_b, x, y)
    {
        scanf("%d %d %d %d", &r_a, &r_b, &x, &y);
    }

    *a = r_a; *b = r_b;
    use_values(id, *a, *b);
}

...
```

OpenMP Syntax

Example (cont'd):

```
int main(int argc, char *argv[] ) {

    int tid = 0;
    int a, b;

    #pragma omp parallel private(tid,a,b)
    {
        #ifdef _OPENMP
            tid = omp_get_thread_num();
        #endif

        init(tid, &a, &b);
        printf("In main - TID %d: a = %d, b = %d,",
               " x = %d, y = %d\n",
               tid, a, b, x, y);
    } /* end of omp parallel */

    return(0);
}
```

OpenMP Syntax

Example output:

```
$ OMP_NUM_THREADS=3 ./copypriv
1 2 3 4
      TID 2: a = 1, b = 2, c = 3, d = 4
      TID 1: a = 1, b = 2, c = 3, d = 4
      TID 0: a = 1, b = 2, c = 3, d = 4

In main - TID 0: a = 1, b = 2, x = 3, y = 4
In main - TID 1: a = 1, b = 2, x = 3, y = 4
In main - TID 2: a
```

without copyprivate
on r_a and r_b

```
env OMP_NUM_THREADS=3 ./copypriv
1 2 3 4
      TID 1: a = 0, b = 1, c = 3, d = 4
      TID 0: a = 0, b = 0, c = 3, d = 4
      TID 2: a = 1, b = 2, c = 3, d = 4

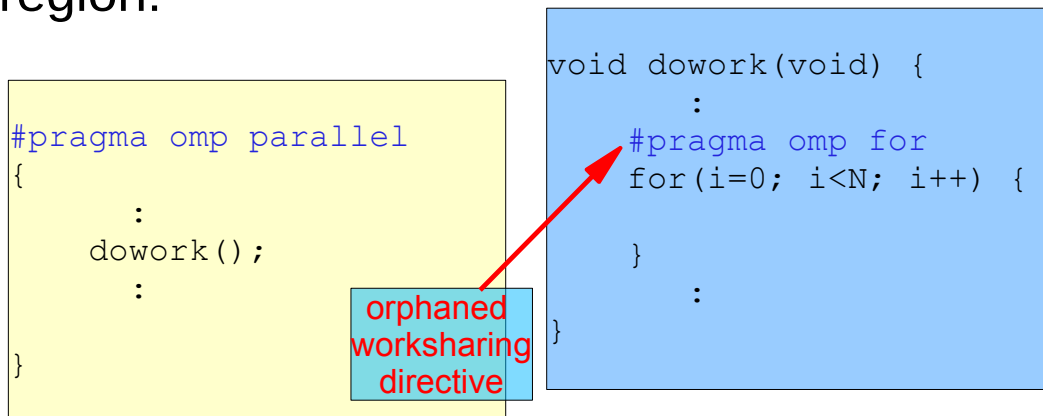
In main - TID 2: a = 1, b = 2, x = 3, y = 4
In main - TID 0: a = 0, b = 0, x = 3, y = 4
In main - TID 1: a = 0, b = 1, x = 3, y = 4
```

OpenMP Orphaning

Orphaning in OpenMP

OpenMP Orphaning

The OpenMP standard does not restrict worksharing and synchronization directives to be within the lexical extent of a parallel region. Those directives can be orphaned, i.e. they can appear outside a parallel region:



OpenMP Orphaning

- When an orphaned directive is detected within the dynamic extent of a parallel region, its behaviour is similar to the non-orphaned case.
- When an orphaned directive is detected in the sequential part of the program, it will be ignored.

```

dowork(); // serial for

#pragma omp parallel
{
    :
    dowork(); // parallel for
    :
}
  
```

```

void dowork(void) {
    :
    #pragma omp for
    for(i=0; i<N; i++) {
        :
    }
}
  
```


Functionality added in OpenMP 3.0

OpenMP syntax

Tasking

- ❑ allows parallelization of work that is generated dynamically
- ❑ provides a flexible model for irregular parallelism
- ❑ uses a “task pool” concept
- ❑ new opportunities:
 - ❑ while loops
 - ❑ recursive structures

OpenMP syntax

❑ Syntax C/C++:

```
#pragma omp task [clause]
{
    ...
}
```

❑ clause can be

- ❑ if (int_expr)
- ❑ default(shared|none)
- ❑ private(list), shared(list)
- ❑ firstprivate(list)
- ❑ untied

OpenMP syntax

❑ Syntax Fortran:

```
!$OMP task [clause]
...
!$OMP end task
```

❑ where clause can be

- ❑ if (int_expr)
- ❑ default(shared|private|firstprivate|none)
- ❑ private(list), shared(list)
- ❑ firstprivate(list)
- ❑ untied

OpenMP syntax

Tasking example I:

while loop:

```
p = lhead;
while (p != NULL)
{
    do_work(p);
    p = next(p);
}
```

parallel while loop with OpenMP tasks:

```
#pragma omp parallel
{
    #pragma omp single
    {
        p = lhead;
        while (p != NULL) {
            #pragma omp task
            {
                do_work(p);
            }
            p = next(p);
        }
    } // end of single
} // end of parallel
```

OpenMP syntax

What's going on?

```
#pragma omp parallel
{
    #pragma omp single
    {
        p = lhead;
        while (p != NULL) {
            #pragma omp task
            {
                do_work(p);
            }
            p = next(p);
        }
    } // end of single
} // end of parallel
```

← start of parallel region

← one thread only, please

← task generation – tasks are added to the task list

← **all work is done here!**

← implicit barrier – all unfinished tasks have to be finished

OpenMP syntax

- ❑ Tasks and recursion: calculating Fibonacci numbers
- ❑ Recursive scheme to calculate the n^{th} Fibonacci number:
 - ❑ $\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$
 - ❑ stopping criterion: return 1 if $n < 2$
- ❑ Caveat: this method is not very effective, but used here to demonstrate the concept of tasking!

OpenMP syntax

The sequential code:

```
int
main(int argc, char* argv[]) {
    [...]
    fib(input);
    [...]
}
```

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;

    x = fib(n - 1);
    y = fib(n - 2);

    return(x + y);
}
```

OpenMP syntax

OpenMP version of fib() with tasks:

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;

    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);

    #pragma omp taskwait
    return(x + y);
}
```

note the special
scoping rules!

generate two tasks,
calling fib() recursively

task synchronization -
to get the right results

OpenMP syntax

Scoping rules with tasks:

- ❑ Static and global variables are shared
- ❑ Local (aka automatic) variables are private
- ❑ Orphaned task variables are firstprivate
- ❑ Non-orphaned task variables inherit the shared attribute
- ❑ (Local) Task variables are firstprivate, unless declared shared
- ❑ Thus, we have to declare x and y as shared

OpenMP syntax

Task synchronization:

- ❑ `#pragma omp taskwait`
- ❑ suspends the encountering task, until all child tasks are completed
- ❑ direct children only, not descendants
- ❑ needed here, to make sure that x and y are still exist when we take the sum.



OpenMP syntax

OpenMP version of main() with tasks:

```
int
main(int argc, char* argv[]) {
    [...]

    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    } // end of omp parallel

    [...]
}
```

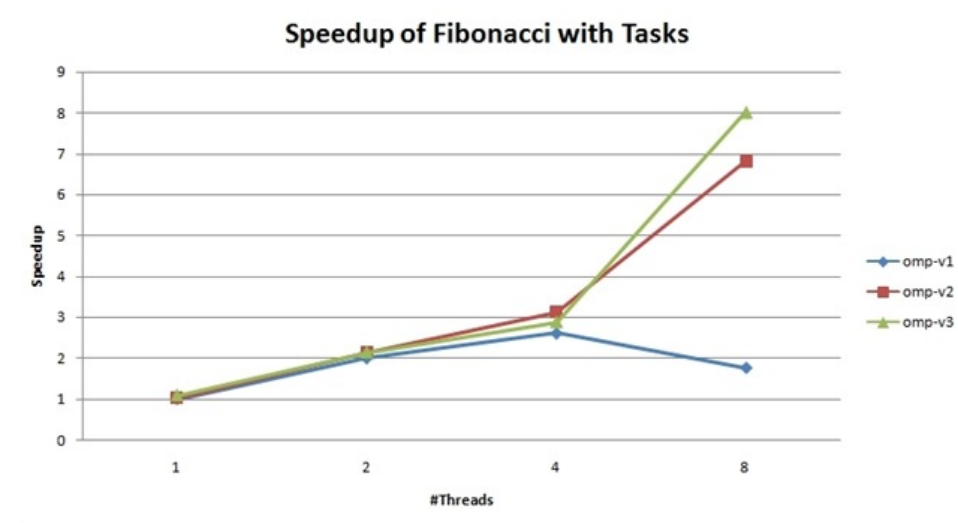
start of parallel region -
team of worker threads

task generation by one
thread, only!



OpenMP syntax

Results of the Fibonacci program



courtesy: Chr. Terboven, RWTH Aachen

OpenMP syntax

Notes on the Fibonacci speedup results:

- ❑ The simple OpenMP version (omp-v1) doesn't scale – as expected – due to the large amount of tasks generated
- ❑ Improvement 1 (omp-v2):
 - ❑ add an if-clause to the tasks:
`#pragma omp task if(n>=30) shared(...)`
 - ❑ improves the speed-up, but still not perfect
- ❑ Improvement 2 (omp-v3): (see next slide)

OpenMP syntax

version omp-v3 of fib() with tasks:

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;
    if (n < 30) {
        return(fib(n-1) + fib(n-2));
    }

    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);

    #pragma omp taskwait
    return(x + y);
}
```

OpenMP: Error detection

Tools to check your OpenMP code

OpenMP compile-time checks

OpenMP checks supported by Oracle Studio compilers:

- ❑ -xloopinfo – info on loops
- ❑ -xvpara – compile time warnings on ...
 - ❑ ... scoping problems
 - ❑ ... possible data races

OpenMP run-time checks

OpenMP run-time checks supported by Oracle Studio:

- ❑ `SUNW_MP_WARN = [true|false]`
 - ❑ gives you warnings at runtime
 - ❑ e.g. dynamic change of threads
 - ❑ inconsistencies of barriers
 - ❑ ...
- ❑ Note: there is a certain performance penalty

OpenMP run-time checks

Example: illegal usage of a barrier

```

1 #include <omp.h>
2 #include <stdio.h>
3
4 int main(void)
5 {
6     #pragma omp parallel num_threads(4)
7     {
8         int i = omp_get_thread_num();
9
10        if (i % 2) {
11            printf("At barrier 1.\n");
12            #pragma omp barrier
13        }
14    }
15    return 0;
16 }

```

OpenMP run-time checks

Example (cont'd): illegal usage of a barrier

```

$ cc -g -xO3 -xopenmp -xvpara -o bad1 bad1.c
$ ./bad1
At barrier 1.
At barrier 1.
^C
$ SUNW_MP_WARN=true ./bad1
WARNING (libmtask): Environment variable SUNW_MP_WARN is set
to TRUE. Runtime error checking will be enabled.
At barrier 1.
At barrier 1.
WARNING (libmtask): Threads at barrier from different directives.
Thread at barrier from bad1.c:6.
Thread at barrier from bad1.c:12.
Possible Reasons:
Worksharing constructs not encountered by all threads in
the team in the same order.
Incorrect placement of barrier directives.
WARNING (libmtask): Runtime shutting down while some parallel
region is still active.
$

```

no compiler warning!!

OpenMP: Data Race Detection

Solaris Studio has a data race detection tool:

- ❑ Thread Analyzer (tha)
- ❑ Quick usage guide:
 - ❑ compile and link with `-xinstrument=datarace`
 - ❑ run it: `collect -r on a.out`
 - ❑ view results (GUI): `tha tha.1.er`
 - ❑ or CLI: `er_print -races tha.1.er`

OpenMP: Data Race Detection

Example:

```
int main(int argc, char *argv[]) {  
  
    int i, total = 0, N = 2000000;  
    int primes[N];  
    #pragma omp parallel for  
    for( i = 2; i < N; i++ ) {  
        if ( is_prime(i) ) {  
            primes[total] = i;  
            total++;  
        }  
    }  
  
    printf("# of prime numbers between 2 and %d: %d\n",  
          N, total);  
    return(0);  
}
```

OpenMP: Data Race Detection

Example (cont'd): compile and run

```
$ cc -g -fast -o prime prime.c
$ ptime ./prime
# of prime numbers between 2 and 2000000: 148933

real      10.862
user      10.483
sys       0.056

$ cc -g -fast -xopenmp -xloopinfo -o prime prime.c
$ "prime.c", line 7: PARALLELIZED, user pragma used
$ OMP_NUM_THREADS=4 ptime ./prime
# of prime numbers between 2 and 2000000: 148310

real      3.662    <--- speed-up: 2.9x
user      10.494
sys       0.055
```

OpenMP: Data Race Detection

Example (cont'd): run, run, ... and use collect

```
$ OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 148310

$ OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 148328

$ cc -g -fast -xopenmp -xinstrument=datarace \
  -o prime prime.c

$ OMP_NUM_THREADS=4 collect -r on ./prime
Creating experiment database tha.1.er ...
# of prime numbers between 2 and 2000000: 125581

$
```

OpenMP: Data Race Detection

Example (cont'd): analyze the collect data

```
$ er_print -races tha.1.er
Total Races: 2 Experiment: tha.1.er

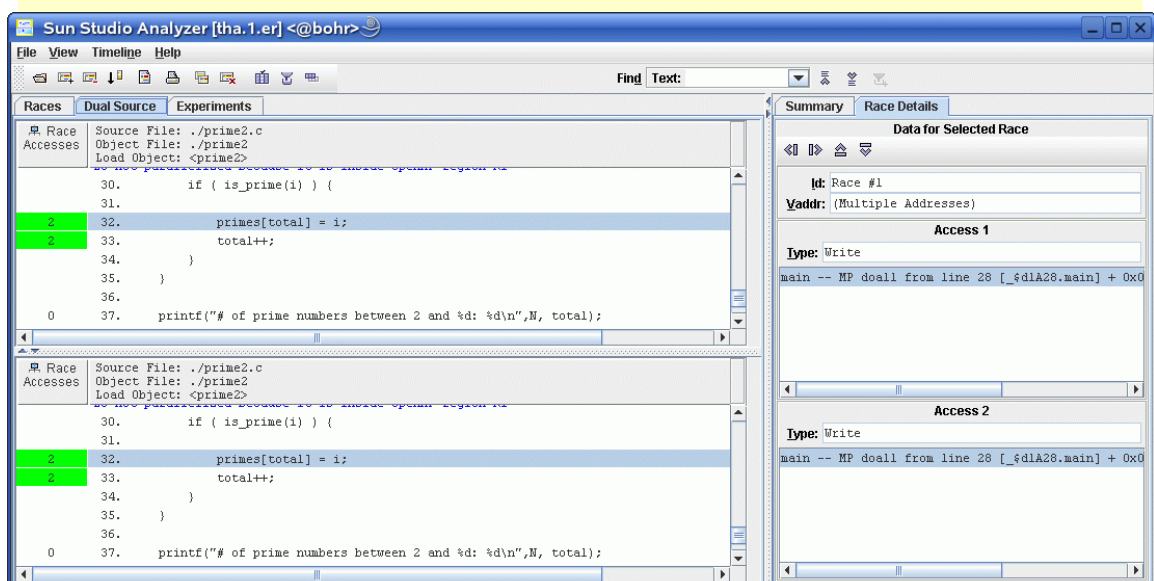
Race #1, Vaddr: (Multiple Addresses)
    Access 1: Write, main -- MP doall from line 6 ...
               line 10 in "prime.c"
    Access 2: Write, main -- MP doall from line 6 ...
               line 10 in "prime.c"
Total Traces: 1

Race #2, Vaddr: 0xffbfecb4
    Access 1: Write, main -- MP doall from line 6 ...
               line 11 in "prime.c"
    Access 2: Write, main -- MP doall from line 6 ...
               line 11 in "prime.c"
Total Traces: 1
```

OpenMP: Data Race Detection

Example (cont'd): analyze the collect data

```
$ tha tha.1.er
```



OpenMP: Data Race Detection

Example (cont'd): fix the bug

```
int main(int argc, char *argv[]) {

    int i, total = 0, N = 2000000;
    int primes[N];
    #pragma omp parallel for

    for( i = 2; i < N; i++ ) {
        if ( is_prime(i) ) {
            #pragma omp critical
            { primes[total] = i;
              total++;
            }
        }
    }
    printf("# of prime numbers between 2 and %d: %d\n",
          N, total);
    return(0);
}
```

January 2018

02614 - High-Performance Computing

47

OpenMP: Data Race Detection

Example (cont'd): check – and recompile

```
$ cc -g -fast -xopenmp -xinstrument=datarace \
  -o prime prime.c

$ OMP_NUM_THREADS=4 collect -r on ./prime
Creating experiment database tha.2.er ...
# of prime numbers between 2 and 2000000: 148933

$ er_print -races tha.2.er
Total Races: 0 Experiment: tha.2.er

$ cc -g -fast -xopenmp -o prime prime.c
$ OMP_NUM_THREADS=4 ptime ./prime
# of prime numbers between 2 and 2000000: 148933

real      3.561
user      10.393
sys       0.051
```

January 2018

02614 - High-Performance Computing

48

OpenMP Scheduling

Controlling the scheduling of OpenMP threads



OpenMP Scheduling

Load balancing:

- ❑ Important aspect of performance
- ❑ Especially for less regular workloads, e.g.
 - ❑ transposing a matrix
 - ❑ multiplications of triangular matrices
 - ❑ parallel searches in a linked list
- ❑ The **schedule** clause provides different iteration scheduling algorithms for loops



OpenMP Scheduling

The “schedule” clause:

```
#pragma omp for schedule(static[,chunk])  
#pragma omp for schedule(dynamic[,chunk])  
#pragma omp for schedule(guided[,chunk])  
#pragma omp for schedule(auto) - new in 3.0  
#pragma omp for schedule(runtime)
```

- ❑ If there is no schedule clause, the default is static.



OpenMP Scheduling

```
#pragma omp for schedule(static[,chunk])
```

Static schedule:

- ❑ Iterations are divided into pieces of size chunk and then **statically** assigned to the threads.
- ❑ If chunk is not defined, the work (N) is equally divided among the number of threads (P), i.e. $\text{chunk} = N/P$.



OpenMP Scheduling

```
#pragma omp for schedule(dynamic[,chunk])
```

Dynamic schedule:

- ❑ Iterations are divided into pieces of size chunk and then **dynamically** assigned to the threads – i.e. when a thread has finished one chunk, it is assigned a new one.
- ❑ The default chunk size is 1.



OpenMP Scheduling

```
#pragma omp for schedule(guided[,chunk])
```

Guided schedule:


- ❑ The chunk size is exponentially reduced with each chunk that gets **dynamically** assigned to the threads; chunk defines the minimum number of iterations to assign each time.

$$\text{chunk} = \text{unass_iter} / (\text{weight} * \text{n_thr})$$

- ❑ The default minimum chunk size is 1.



OpenMP Scheduling



OpenMP Scheduling

```
#pragma omp for schedule(runtime)
```

Runtime schedule:

- ❑ The schedule is detected at runtime from the setting of the OMP_SCHEDULE environment variable.
- ❑ Syntax: OMP_SCHEDULE=type,chunk

OpenMP Scheduling

```
#pragma omp for schedule(auto)
```


Automatic schedule (new in OpenMP 3.0):

- When `schedule(auto)` is specified, the decision regarding scheduling is delegated to the compiler and/or runtime system. The programmer gives the implementation the freedom to choose any possible mapping of iterations to threads in the team.

**A real world example:
Molecular Dynamics simulation**

Example: MD simulation

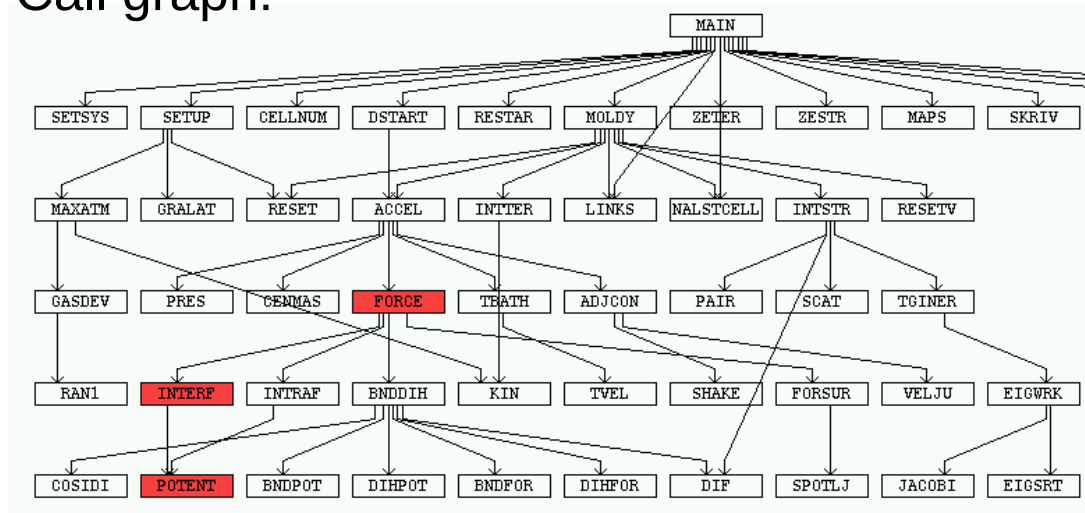
- ❑ Molecular Dynamics simulation of long carbon molecules on a surface:



- ❑ 7200+ lines of Fortran 77 code
- ❑ GOTOs, COMMON blocks, ...
- ❑ one source file

Example: MD simulation

Call graph:



more than 80% of the runtime are spent in the red part of the call graph

Example: MD simulation

- ❑ The loop to be parallelized contains a call to another subroutine.
- ❑ Data is passed the old Fortran style via COMMON blocks
- ❑ First try: Inserted one PARALLEL DO pragma in the code, using autoscoping, i.e. a feature of the Solaris Studio compiler (more later)
- ❑ The compiler generated a parallel version!

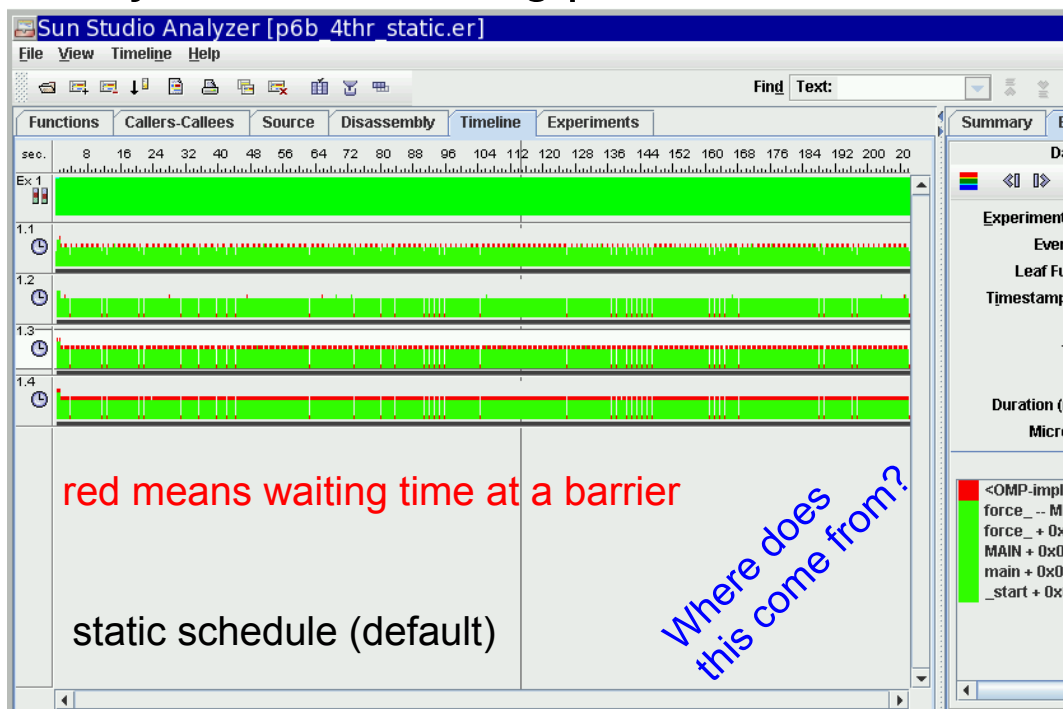
This took us by surprise!

Example: MD simulation

- ❑ First test runs:
 - ❑ It didn't scale ...
 - ❑ The results were dependent on the number of threads ...
- ❑ Thread analyzer revealed data races in two variables inside the called subroutine.
- ❑ Fix: Added additional scoping for those variables in the OpenMP pragma!
- ❑ This solved the data race problem.

Example: MD simulation

Analysis of the scaling problem:



January 2018


02614 - High-Performance Computing

63

Example: MD simulation

The simulation box:

seen from the side



← thread 4

← thread 3

← thread 2

← thread 1

subdivision into smaller cells

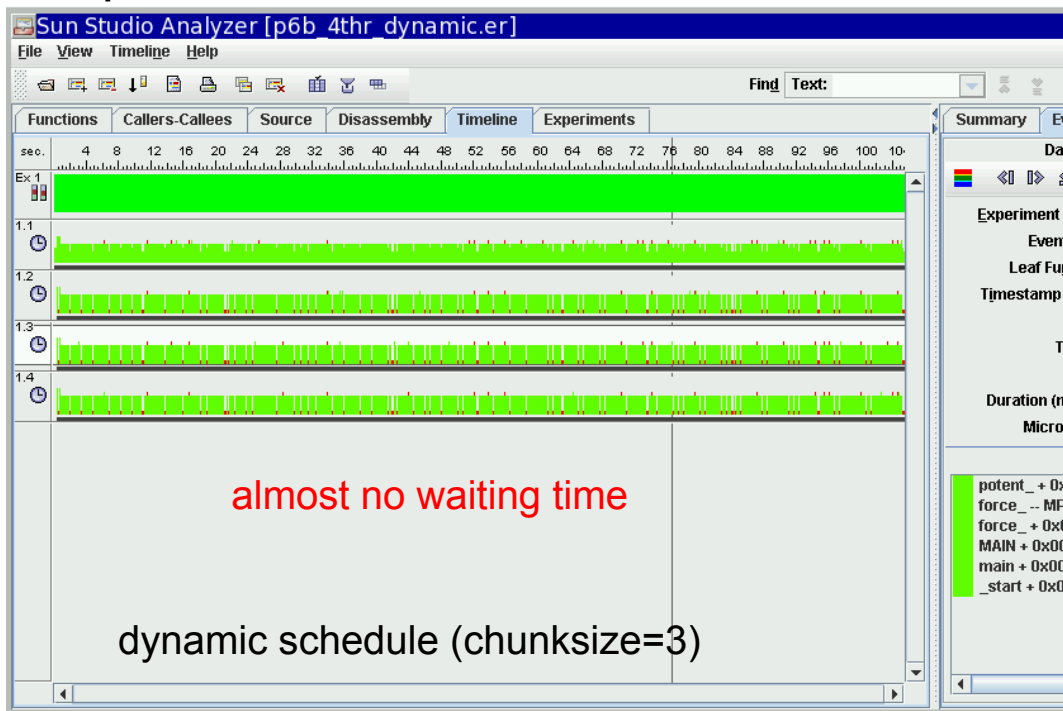
January 2018

02614 - High-Performance Computing

64

Example: MD simulation

Adapted the schedule:




January 2018

02614 - High-Performance Computing

65

Example: MD simulation

Speed-up results:



January 2018

02614 - High-Performance Computing

66