

Parallel Programming in OpenMP – part III

Outline

- ❑ Runtime library
- ❑ Environment variables
- ❑ OpenMP Future
- ❑ Behind the scenes
- ❑ Summary
- ❑ References

OpenMP Runtime Library

The OpenMP runtime library:
support functions



OpenMP Runtime Library

The OpenMP standard defines an API for library calls, that have a variety of functions:

- ❑ query
 - ❑ the number of threads/processors
 - ❑ thread ID, “in parallel”
- ❑ set
 - ❑ the number of threads to use
 - ❑ scheduling mode
- ❑ locking (semaphores)



OpenMP Runtime Library

Name

omp_set_num_threads
omp_get_num_threads
omp_get_max_threads
omp_get_thread_num
omp_get_num_procs
omp_in_parallel

omp_set_dynamic
omp_get_dynamic

omp_set_nested
omp_get_nested

omp_get_wtime
omp_get_wtick

Functionality

set number of threads
get number of threads in team
get max. number of threads
get thread ID
get max. number of processors
check whether in parallel region

activate dynamic thread adjustment
check for dynamic thread adjustment
(implementation can ignore this)

activate nested parallelism
check for nested parallelism
(implementation can ignore this)

returns wall clock time
number of second between clock ticks



OpenMP Runtime Library

Function prototypes:

```
void omp_set_num_threads(int num_threads)
int  omp_get_num_threads(void)
int  omp_get_max_threads(void)
int  omp_get_thread_num(void)
int  omp_get_num_procs(void)
int  omp_in_parallel(void)
```

```
void omp_set_dynamic(int dynamic_threads)
int  omp_get_dynamic(void)
void omp_set_nested(int nested)
int  omp_get_nested(void)
```

```
double omp_get_wtime(void)
double omp_get_wtick(void)
```



OpenMP 3.0 Runtime Library

Name

`omp_set_schedule`
`omp_get_schedule`

Functionality

set the schedule
 get the schedule

`omp_get_thread_limit`

max. number of available threads
 in the implementation

`omp_set_max_active_levels`
`omp_get_max_active_levels`
`omp_get_level`
`omp_get_ancestor_thread_num`

set the number of nested levels
 get the number of nested levels
 returns the current nesting level
 returns thread id of the ancestor
 thread in specified level

`omp_get_team_size`
`omp_get_active_level`

get team size at specified level
 returns the number of enclosing,
 active nested parallel regions

for more details see the OpenMP 3.0 specifications

OpenMP Runtime Library

Usage of `omp_get_num_threads()` vs `omp_get_max_threads()`:

```
// get the number of threads
threads = omp_get_max_threads();
```

returns value of `OMP_NUM_THREADS`

```
// get the number of threads
threads = omp_get_num_threads();

#pragma omp parallel
{
  #pragma omp master
  { threads = omp_get_num_threads(); }
} // end parallel
```

returns 1- outside a parallel region

returns value of threads in a parallel region

OpenMP Runtime Library

Measuring time:

- ❑ It is most useful to compare wall clock times

```
double ts, te;  
ts = omp_get_wtime();  
  
do_work();  
  
te = omp_get_wtime() - ts;  
  
printf("Elapsed time: %lf\n", te);
```

- ❑ clock() returns the accumulated CPU time of all threads!

OpenMP Environment Variables

Controlling OpenMP via Environment Variables

OpenMP Environment Variables

- ❑ `OMP_NUM_THREADS = n`
 - ❑ sets the max. no of threads to n (default: 2)
- ❑ `OMP_SCHEDULE = schedule[,chunk]`
 - ❑ schedule: [static | guided | dynamic]
 - ❑ chunk: size of chunks (defaults: [n/a|1|1])
 - ❑ Note: applies to parallel do/for loops only!
- ❑ `OMP_DYNAMIC = [TRUE | FALSE]`
- ❑ `OMP_NESTED = [TRUE | FALSE]`



OpenMP 3.0 Environment Variables

- ❑ `OMP_STACKSIZE = size[B|K|M|G]`
 - ❑ sets the size of the stack of OpenMP threads
 - ❑ default unit: Kilobytes
- ❑ `OMP_WAIT_POLICY = active|passive`
 - ❑ controls the behaviour of idle threads
 - ❑ active: “spinning threads”, i.e. use cycles
 - ❑ passive: threads go to sleep
 - ❑ the default is implementation dependent



OpenMP 3.0 Environment Variables

- ❑ `OMP_MAX_ACTIVE_LEVELS = n`
 - ❑ controls the max. level for nested parallelism
- ❑ `OMP_THREAD_LIMIT = n`
 - ❑ sets the maximum number of threads for an OpenMP program



OpenMP Environment Variables

Oracle Studio specific variables:

- ❑ `SUNW_MP_WARN = [TRUE | FALSE]`
 - ❑ issues warnings, e.g. when requesting too many threads, ...
- ❑ `SUNW_MP_THR_IDLE = [SPIN | SLEEP(τ)]`
- ❑ behaviour of the idle threads
- ❑ τ is the time (in seconds/milliseconds – default: 5 ms) the idle threads spin before they go to sleep
- ❑ Ex.: `SUNW_MP_THR_IDLE=SLEEP(50ms)`
- ❑ OpenMP 3.0: use `OMP_WAIT_POLICY` !



OpenMP Environment Variables

Notes:

- ❑ All the defaults (**in green**) given above are for the Oracle Studio OpenMP implementation.
- ❑ The max. number of threads is limited to the number of on-line processors (cores) in the system. This can be changed by setting `OMP_DYNAMIC` to `FALSE` – be careful when playing with this.
- ❑ Check with your compiler documentation, what the defaults are for different OpenMP implementations, e.g. Intel, GCC, or

OpenMP Precedence

- ❑ Level of priority:
 - 1 clauses, e.g. `num_threads(...)`
 - 2 library calls, e.g. `omp_set_num_threads(...)`
 - 3 environment variables, e.g. `OMP_NUM_THREADS`
- ❑ For a detailed discussion see the OpenMP specifications or check the documentation of your OpenMP implementation.

OpenMP Future

OpenMP standard extensions:
Coming soon to a compiler near you ...

OpenMP Future: Autoscoping

Courtesy: Dieter an Mey, RWTH Aachen

```
!$omp parallel do &
!$omp &
...
!$omp & omegaz,prode,qdens,qjc,qmqc,redbme,redbpe,renbme,&
omp & renbpe, resbme, resbpe, reubme, reubpe, rkdbmk, rkdbpk, rknbm, &
omp & rknbpk, rksbmk, rksbpk, rkubmk, rkubpk, rtdbme, rtdbpe, rtnbme, &
omp & rtnbpe, rtsbme, rtsbpe, rtubme, rtubpe, rudbme, rudbmx, rudbmy, &
omp & rudbmz, rudbpe, rudbpx, rudbpy, rudbpz, runbme, runbmz, runbmy, &
omp & runbmz, runbpe, runbpx, runbpy, runbpz, rusbme, rusbmz, rusbmy, &
omp & rusbmz, rusbpe, rusbpx, rusbpy, rusbpz, ruubme, ruubmx, ruubmy, &
omp & ruubmz, ruubpe, ruubpx, ruubpy, ruubpz, rvdme, rvdmx, rvdmy, &
omp & rvdbmz, rvdpe, rvdpx, rvdpy, rvdpz, rvnme, rvnmx, rvnmy, &
omp & rvnbmz, rvnpe, rvnpx, rvnpy, rvnbpz, rvsbme, rvsbmz, rvsbmy, &
omp & rvsbmz, rvspe, rvspx, rvsby, rvsbpz, rvubme, rvubmx, rvubmy, &
omp & rvubmz, rvubpe, rvubpx, rvubpy, rvubpz, rwdme, rwdmx, rwdmy, &
omp & rwdbmz, rwdpe, rwdpx, rwdpy, rwdpz, rwnme, rwnmx, rwnmy, &
omp & rwnbmz, rwnpe, rwnpx, rwnpy, rwnbpz, rwsbme, rwsbmz, rwsbmy, &
omp & rwsbmz, rwspe, rwspx, rwsby, rwsbpz, rwubme, rwubmx, rwubmy, &
omp & rwubmz, rwubpe, rwubpx, rwubpy, rwubpz, rwdme, rwdmx, rwdmy, &
omp & tdbp, teb, tkc, tk
omp & tknbm, tknbp, tks
omp & tkwb, tnb, tnbm, t
omp & tubm, tubp, twb, u
omp & unb, unbm, unbp, u
omp & uubp, uwb, vc, vdb
omp & vnb, vnbm, vnbp, v
omp & vubp, vwb, wc, wdb
omp & wnbm, wnbp, wsb, wsbm, wsbp, wub, wubm, wubp, &
omp & ww, xiabc, xiabeb, xiabwb, xiayc, xiayeb, xiaywb, xiabc, &
omp & xiabeb, xiabwb, xibxc, xibxnb, xibxsb, xibynb, xibysb, xibznb, &
!$omp & xibzsb, xicxc, xicxdb, xicxub, xicydb, xicyub, xiczdb, xiczub)
do i = is, ie
---- 1600 lines omitted ----
end do
```

OpenMP Future: Autoscopying

- ❑ available with the Oracle Studio compilers
- ❑ if the compiler can't autoscope, you will get a message why it failed
 - ❑ use `-xvpara` to see the messages
 - ❑ the failure message is on the `.o` file as well, make it visible with the `er_src` command
- ❑ is a proposed extension for an upcoming OpenMP standard (didn't make it into OpenMP 3.0, 4.0, 4.5, ...)



OpenMP Future

- ❑ More extensions were/are discussed in the OpenMP ARB and the community, and made or make it into the standard, e.g. extensions for
 - ❑ better performance
 - ❑ memory placement (4.0)
 - ❑ debugging
 - ❑ checks, both at compile- and run-time
 - ❑ exception handling (4.0)
 - ❑ access to accelerators (e.g. GPUs) (4.0)
 - ❑ ...



OpenMP: Behind the scenes

What the compiler does
with your code



OpenMP: Behind the scenes

```
#define MAX_SIZE 8000000
int main() {
    double GlobSum;           /* A global variable */
    double array[MAX_SIZE];
    int nthreads;
    int i;
    /* Initialize things */
    for (i=0; i<MAX_SIZE; i++) array[i] = i;
    GlobSum = 0;
    nthreads = omp_get_max_threads();
    printf("Threads: %d\n", nthreads );
    #pragma omp parallel for private(i) \
        reduction(+ : GlobSum)
    for(i=0; i<MAX_SIZE;i++)
        GlobSum = GlobSum + array[i];

    return(EXIT_SUCCESS);
}
```



OpenMP: Behind the scenes

- ❑ Used the OMPi compiler to generate the intermediate code shown on the next slides.
- ❑ The actual implementation differs from compiler to compiler, and probably also from version to version (improvements).

OpenMP: Behind the scenes

```
int main() {
    ...
    int i;
    _omp_initialize();

    for (i = 0; i < 8000000; i++) array[i] = i;
    GlobSum = 0;
    nthreads = omp_get_max_threads();
    printf("Threads: %d\n", nthreads);

    /* #pragma omp parallel for private(i) reduction(+: GlobSum) */
    {
        _OMP_PARALLEL_DECL_VARSTRUCT(main_parallel_0);
        _OMP_PARALLEL_INIT_VAR(main_parallel_0, GlobSum);
        _OMP_PARALLEL_INIT_VAR(main_parallel_0, array);
        _omp_create_team((-1), _OMP_THREAD, main_parallel_0,
            (void *) &main_parallel_0_var); /* create team of
                                           * threads */
        _omp_destroy_team(_OMP_THREAD->parent);
    }

    return 0;
}
```

OpenMP: Behind the scenes

```
void *main_parallel_0(void *_omp_thread_data){
    int      _omp_dummy = _omp_assign_key(_omp_thread_data);
    double   (*array)[8000000] = &_OMP_VARREF(main_parallel_0,array);
    {
        int      i;
        double   GlobSum = 0;
        int      _omp_start, _omp_end, _omp_incr, _omp_last_iter = 0;
        int      _omp_for_id = _omp_module.for_ofs + 0;
        int      (*_omp_sched_bounds_func) (int, int, int, int,
                                           int, int *, int *, int, int, int *);
        /* static with chunksize or runtime */
        int      _omp_init_start, _omp_nchunks, _omp_c = 0,
                _omp_chunksize;
        _omp_incr = (1);
        _omp_init_directive(_OMP_FOR, _omp_for_id, 0,
                           _omp_incr, 0, 115);
        _omp_sched_bounds_func = _omp_static_bounds;
        _omp_static_bounds_default(8000000, 0, _omp_incr,
                                   &_omp_start, &_omp_end);
        ...
    }
}
```

OpenMP: Behind the scenes

```
...
while ((*_omp_sched_bounds_func) (8000000, 0, _omp_for_id,
    _omp_incr, -1, &_omp_start, &_omp_end, 1, 0, &_omp_c)) {
    if (_omp_start < (8000000) && _omp_end == (8000000))
        _omp_last_iter = 1;

    for (i = _omp_start; i < _omp_end; i++) {
        GlobSum = GlobSum + (*(array))[i];
    }
    /* for */

    if (_omp_last_iter) { /* lastprivate assignments */ }

    /* reduction operation (+:GlobSum) */
    othread_set_lock(&_omp_module.reduction_lock[0]);
    _OMP_VARREF(main_parallel_0, GlobSum) += GlobSum;
    othread_unset_lock(&_omp_module.reduction_lock[0]);
}
return 0;
}
```

OpenMP vs POSIX threads

A possible POSIX threads solution:

```
main() {
    int i,retval;
    pthread_t tid;

    /* Initialize things */
    pthread_attr_init(&attr);
    pthread_mutex_init (&my_mutex, NULL);
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);

    for (i=0; i<MAX_SIZE; i++) array[i] = i;
    GlobSum = 0;

    for(i=0;i<ThreadCount;i++) {
        index[i] = i;
        retval = pthread_create(&tid,&attr,SumFunc,
                                (void *)index[i]);

        thread_id[i] = tid;
    }
    for(i=0;i<ThreadCount;i++)
        retval = pthread_join(thread_id[i],NULL);
}
```

January 2018

02614 - High-Performance Computing

94

OpenMP vs POSIX threads

```
void *SumFunc(void *parm){
    int i,me,chunk,start,end;
    double LocSum;

    /* Decide which iterations belong to me */
    me = (int) parm;
    chunk = MAX_SIZE / ThreadCount;
    start = me * chunk;
    end = start + chunk; /* C-Style - actual element + 1 */
    if ( me == (ThreadCount-1) ) end = MAX_SIZE;

    /* Compute sum of our subset*/
    LocSum = 0;
    for(i=start;i<end;i++ ) LocSum = LocSum + array[i];

    /* Update the global sum and return */
    pthread_mutex_lock (&my_mutex);
    GlobSum = GlobSum + LocSum;
    pthread_mutex_unlock (&my_mutex);
}
```

Note: Variable definitions are omitted in this example!

January 2018

02614 - High-Performance Computing

95

OpenMP Summary

Short summary
of the three lectures



OpenMP Summary

- ❑ OpenMP: a parallel programming model for SMP computers
- ❑ compiler directives, support functions, environment variables
- ❑ easy to implement, also “little by little”
- ❑ next lecture: “OpenMP & Performance”



OpenMP References

- ❑ Useful Websites:

- ❑ <http://www.openmp.org/>
- ❑ <http://www.compunity.org/>

- ❑ Tutorials:

- ❑ <https://computing.llnl.gov/tutorials/openMP/>
- ❑ http://ircc.fiu.edu/download/sc13/AdvOpenMP_Slides.pdf

- ❑ Implementations: search for OpenMP

- ❑ Oracle: search on <http://docs.oracle.com/en/>
- ❑ Intel: search on <http://software.intel.com/>
- ❑ GCC: <https://gcc.gnu.org/wiki/openmp>