

Introduction to Assignment 3

- GPU Matrix Multiplication
 - GPU Poisson Problem
- + Compare with your best CPU versions
- Background:
 - Re-read the two previous assignments
 - Profiler: `nvvp`
 - Brief demonstration at the end of this introduction

GPU Matrix multiplication

GPU Matrix multiplication

■ The `matmult_f.nvcc` driver is provided

```
matmult_f.nvcc type m n k [bs]
```

where `m`, `n`, `k` are the parameters defining the matrix sizes, `bs` is the optional blocksize for the block version, and `type` can be one of:

```
nat      - the native/naive version
lib      - the library version (note that this now calls a multithreaded
library)
gpu1     - the first gpu version
gpu2     - the second gpu version
gpu3     - the third gpu version
gpu4     - the fourth gpu version
gpu5     - the fifth gpu version
gpu6     - the sixth gpu version
gpulib   - the CUBLAS library version
```

as well as `blk`, `mnk`, `nmk`, ... (the permutations).

■ See README for more (also week 1 README)

GPU Matrix multiplication

- Reference version: BLAS (e.g., cblas)

```
void DGEMM(char *transa, char *transb,  
           int *m, int *n, int *k,  
           double *alpha,  
           double *A, int *lda,  
           double *B, int *ldb,  
           double *beta,  
           double *C, int *ldc);
```

- You need to use `extern "C" {}` when including header files for C libraries in `.cu` files

```
extern "C" { #include <cblas.h> }
```

GPU Matrix multiplication

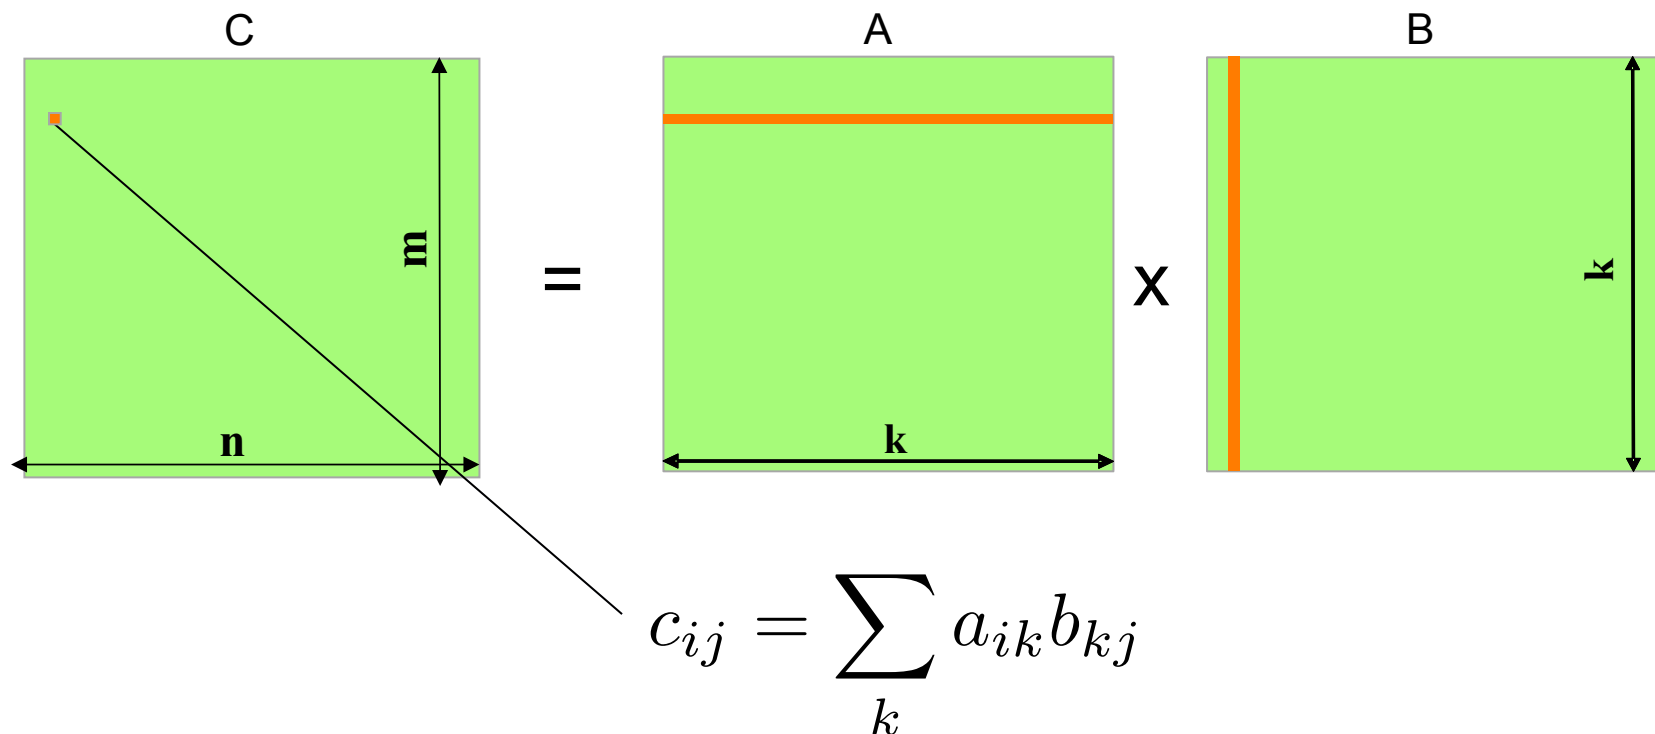
- You also need to use `extern "C" {}` for the functions in your shared library (driver is by `gcc`)

```
extern "C" {  
    matmult_lib(...)  
    {  
        ...  
    }  
    ...etc.  
}
```

- C code in separate `.c` files may be compiled by `gcc` or `nvcc` but always linked in by `nvcc`

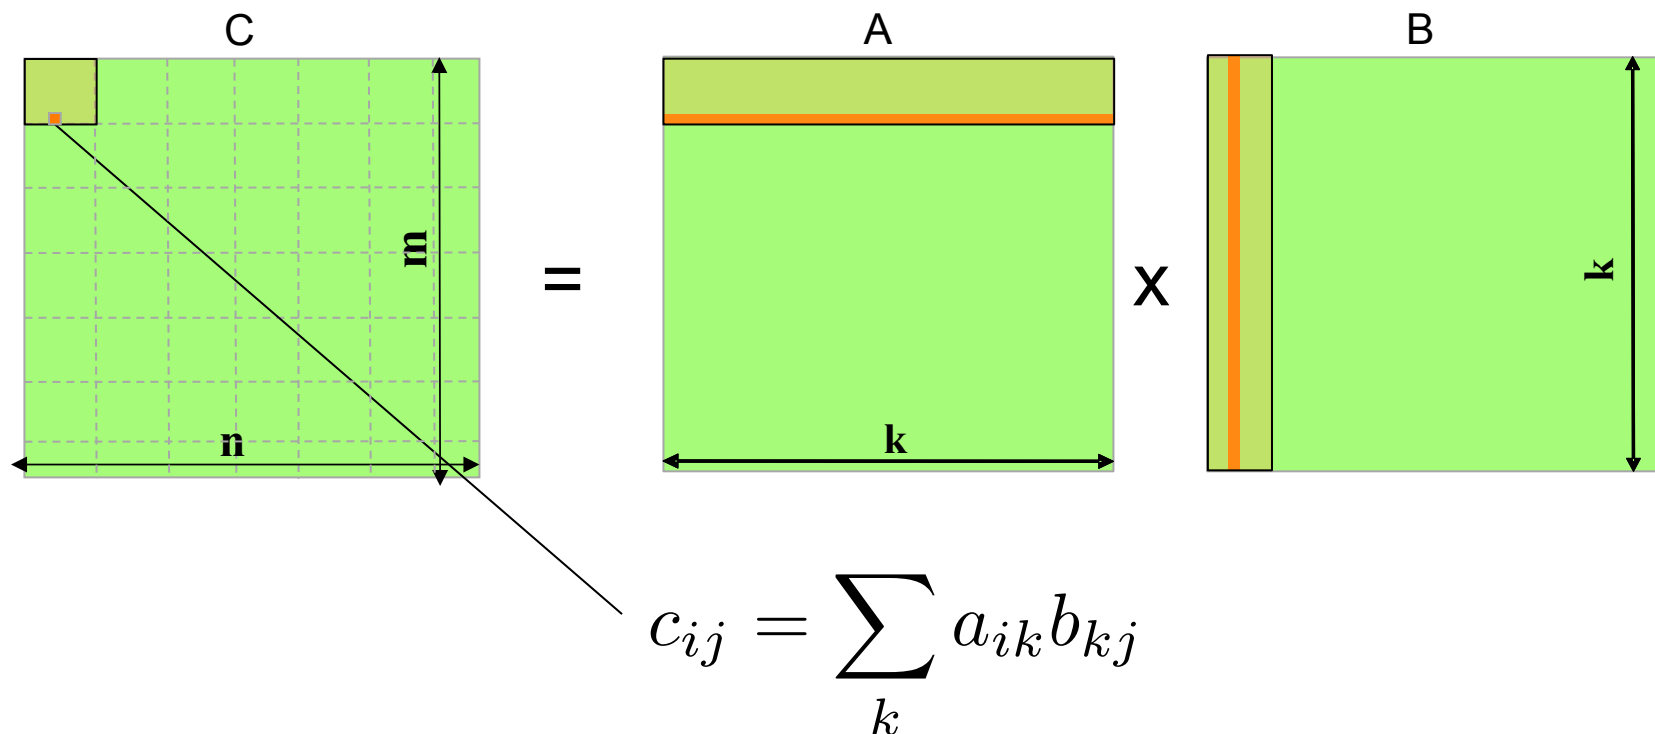
GPU Matrix multiplication

- Sequential version v1: One thread does it all
 - Launch configuration $\langle\langle\langle 1, 1, \rangle\rangle\rangle$



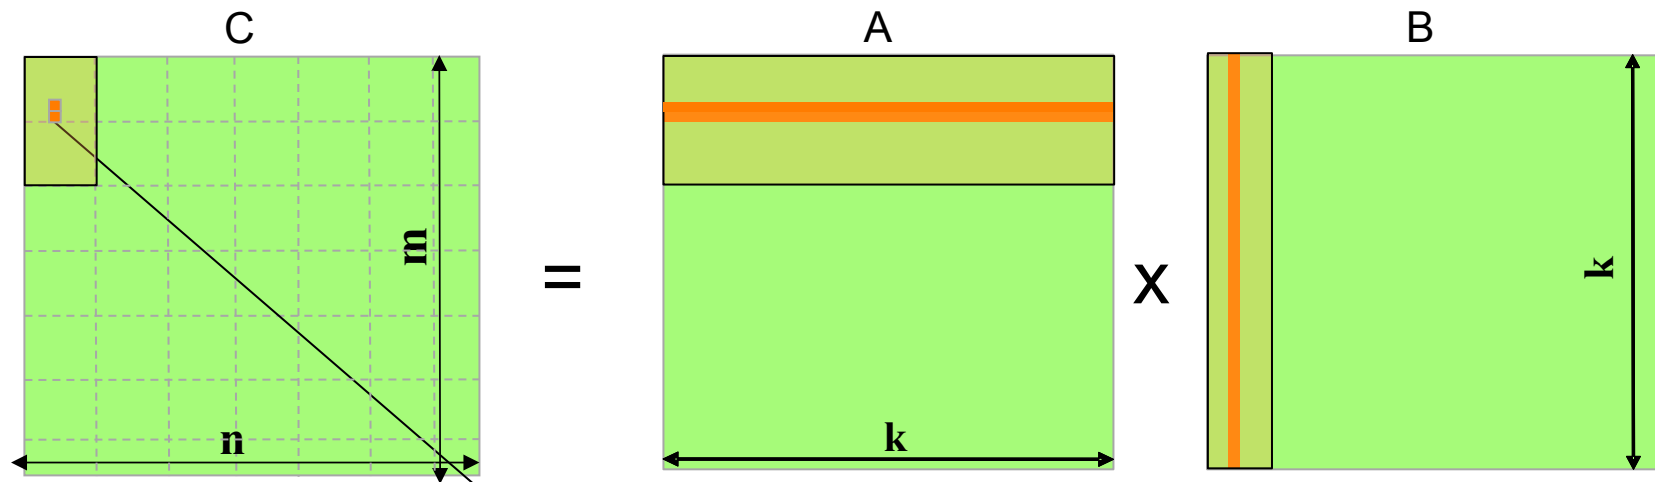
GPU Matrix multiplication

- Naive version v2: One thread per element in C
 - 2D Grid, 2D block (for example 16 x 16 threads)



GPU Matrix multiplication

- Register blocking v3: Each thread does 2 elements
 - Which position of second element is optimal?
 - 2D grid is 1/2 in y dimension, 2D block is the same



$$c_{ij} = \sum_k a_{ik} b_{kj}$$

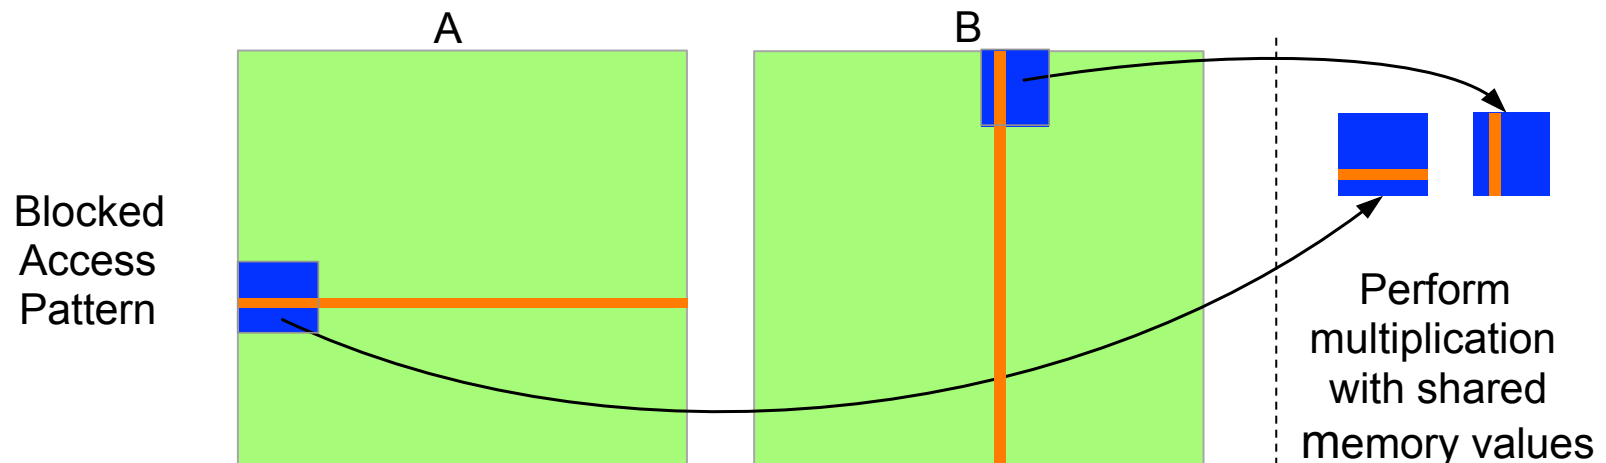
GPU Matrix multiplication

- Extended register blocking v4:
 - Each thread does more than two elements
 - How many are optimal? Which positions are optimal?
 - Easiest way to allocate more elements in registers
 - `double C_reg[4] = {0.0, 0.0, 0.0, 0.0};`
 - `double C_reg[2][2] = {0.0, 0.0, 0.0, 0.0};`
 - Write several trial cases and choose the best
 - Or write a generic kernel that can run all choices, and select the best parameters

GPU Matrix multiplication

■ Shared memory v5: Read in blocks of A of B

- ❑ E.g. use `dim3(16, 16)` blocks and split the 'k' loop in pieces of 16
- ❑ Allocate shared memory: `A_s[16][16]` and `B_s[16][16]`.



■ Go to the Nvidia online documentation:

[http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html - shared-memory](http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html-shared-memory)

- ❑ Copy-paste code from here and modify to fit the matmult driver

GPU Matrix multiplication

■ Compare with cuBLAS version

```
cublasStatus_t cublasDgemm(cublasHandle_t handle,  
                           cublasOperation_t transa,  
                           cublasOperation_t transb,  
                           int m, int n, int k,  
                           const double *alpha,  
                           const double *A, int lda,  
                           const double *B, int ldb,  
                           const double *beta,  
                           double *C, int ldc)
```

- ❑ Note: cuBLAS library uses column-major storage
- ❑ New and legacy cuBLAS APIs (see [here](#))
- ❑ `#include <cublas_v2.h>`

GPU Matrix multiplication

- The `-G` flag sets debug lines into your code
 - This reduces the performance drastically
 - Remove it for performance tuning!!!
- For large matrices please use
 - `MFLOPS_MAX_IT=1 ./matmult_f.nvcc ...`
- For benchmarks please use
 - `MATMULT_COMPARE=0`
- Benchmark runs should be submitted as batch jobs (see last slide)

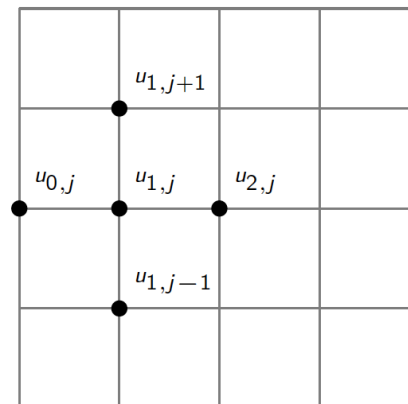
GPU Poisson Problem

GPU Poisson problem

- Reference version: Your best OpenMP version from assignment 2.
 - Also use your code to allocate and initialize the necessary matrices for the square room problem
 - Note that if you used the `cc` sun compiler before there might be slight differences to the `gcc` compiler

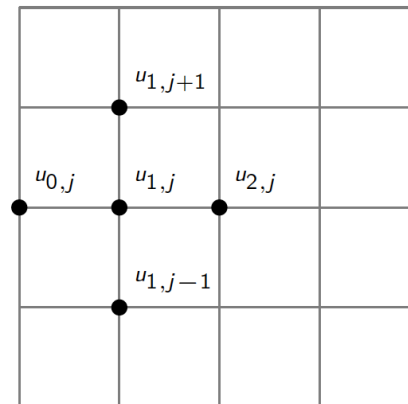
GPU Poisson problem

- Sequential version: One thread does it all
 - Launch configuration $\langle\langle\langle 1, 1, \rangle\rangle\rangle$
 - Hint: Do only one iteration per kernel launch!
 - Hint: Swap pointers for u and u_old on the CPU



GPU Poisson problem

- Naive version: One thread per grid-point
 - 2D grid, 2D block
 - Global memory usage only – rely on caches to help

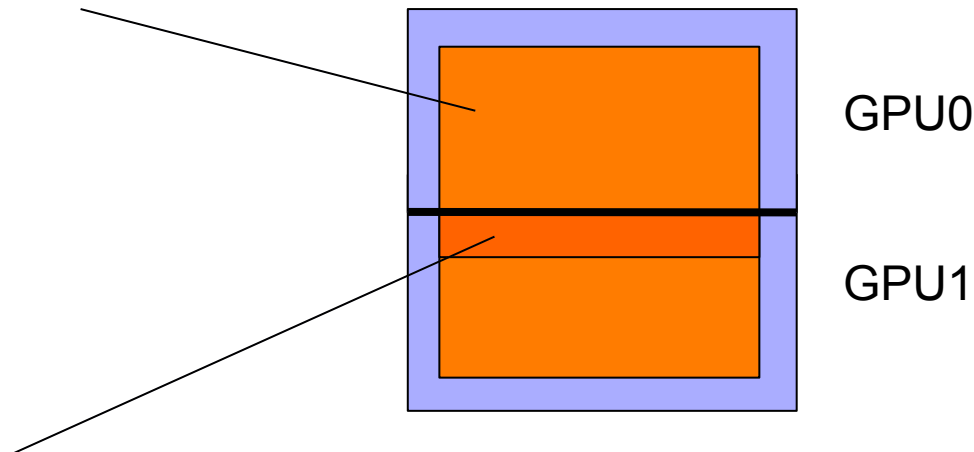


GPU Poisson problem

■ Multi-GPU version

- ❑ Split task into two – top and bottom
- ❑ Interior points can be updated from global memory
- ❑ Border points must read "peer values" from other GPU

Read from global memory



Available from other GPU

Reports / Analysis

- Assess your performance (Gflops / Bandwidth)
- Speed-up calculations (fair)
- Tuning considerations
- Profiler analysis
- Relevant comments and observations

Submitting GPU batch jobs

- Benchmark runs should always be submitted
 - Select queue and ask for GPUs using these options
 - `-q gpuv100`
 - `-gpu "num=1:mode=exclusive_process:mps=yes"`
 - Put these into your job scripts
- Maximum wall-clock time on jobs 1 hour!
- For jobs using two GPUs use `num=2`
- For CPU-only jobs please do not request GPUs